

# **Scalability Analysis: Naive Implementation**

## **Precision Impact & Benchmark vs OpenBLAS**

Project AMSC

November 27, 2025

# Performance Overview: Naive (Float)

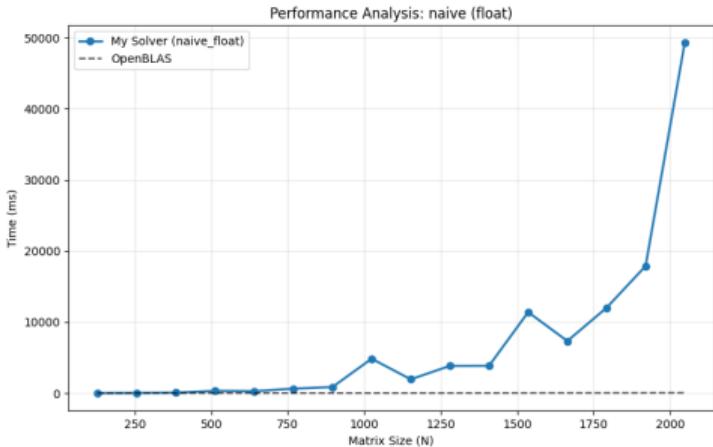


Figure: Execution Time (ms)

*Exponential growth due to cache thrashing.*

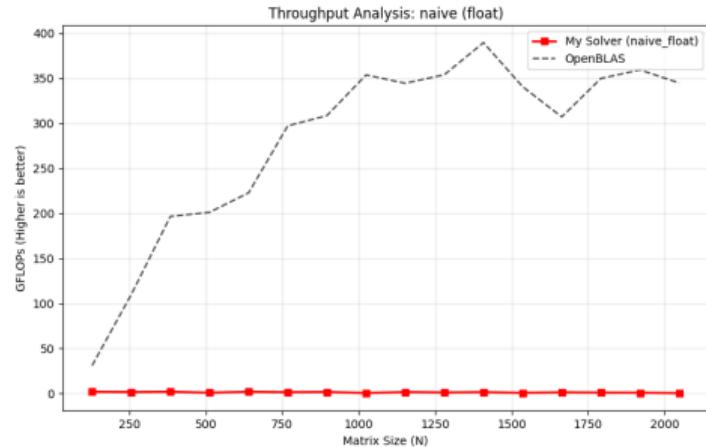


Figure: Throughput (GFLOPS)

*Performance remains  $\approx 2$  GFLOPS regardless of size.*

# Quantitative Analysis: Float vs Double

## Impact of Precision on Execution Time

Size ( $N$ )	Float	Double	$\Delta$ Overhead
$128^3$	2.59 ms	3.85 ms	+48%
$512^3$	278 ms	377 ms	+35%
$1024^3$	3.63 s	5.22 s	+43%
$2048^3$	47.60 s	48.94 s	Only +2.8%

### Insight: The "Latency Wall"

For smaller sizes, double is slower due to higher bandwidth usage (8 bytes vs 4 bytes).

However, at  $N=2048$ , the difference disappears.

**Why?** The CPU is entirely stalled by *latency* (waiting for RAM), not *bandwidth* (transfer speed). The penalty for jumping around in memory masks the data size difference.

Reference: OpenBLAS (Float) at  $N = 2048$  takes only **0.058 s** ( $\approx 820x$  faster).

# Why does Naive fail? (Bottleneck Analysis)

The standard  $(ijk)$  implementation exposes critical hardware limitations:

- **Cache Thrashing (Major):** Matrix  $B$  is accessed column-wise ( $stride = N$ ). This breaks spatial locality, causing a massive amount of L1/L2 cache misses.
- **Precision Penalty:** Using double halves the effective SIMD width (2 doubles vs 4 floats per 128-bit vector) and consumes 2x cache lines, exacerbating the cache pressure.
- **Pipeline Stalls:** The CPU spends most cycles waiting for data fetch rather than computing.

## Memory Access Pattern

$B[k][j]$  with varying  $k$



Jumping  $N \times 4$  bytes every step leads to RAM fetch instead of Cache hit.