

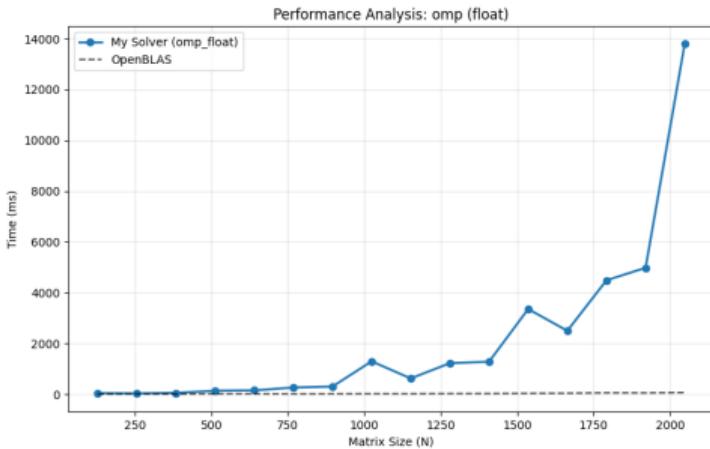
# **Scalability Analysis: OpenMP Implementation**

## **Multithreading Overhead & Bandwidth Saturation**

Progetto AMSC

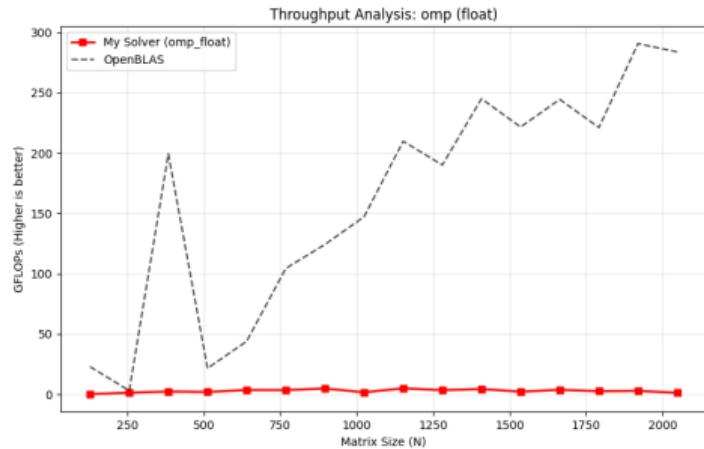
November 27, 2025

# Performance Overview: OpenMP (Float)



**Figure: Execution Time (ms)**

*Significant instability (spikes) due to thread synchronization overhead.*



**Figure: Throughput (GFLOPS)**

*Throughput is unstable and surprisingly lower than SIMD.*

# Quantitative Analysis: Float vs Double

## Impact of Precision on Multithreading

Size ( $N$ )	Float	Double	$\Delta$ Overhead
$128^3$	28.7 ms	25.7 ms	-10% (Noise)
$1024^3$	1.24 s	1.48 s	+19%
$1536^3$	3.51 s	5.06 s	Huge Spike
$2048^3$	14.39 s	16.27 s	+13%

## The "Parallelizing Garbage" Trap

Compared to Naive (47s), OMP is faster (14s,  $\approx 3.3$ x speedup).

**BUT it is slower than SIMD (1.07s)!**

Why? We are parallelizing the inefficient *Naive* logic. Running 8 threads that all experience Cache Misses just saturates the Memory Bus 8 times faster.

**Key Takeaway:** More cores cannot fix bad memory access patterns.

# Why is OpenMP slower than SIMD?

Simply adding `#pragma omp parallel for` does not magically solve architecture bottlenecks:

- **Memory Bandwidth Saturation:** Since the code is not blocked (tiled), every thread constantly fetches data from RAM. The shared memory bus becomes a traffic jam.
- **False Sharing / Cache Contention:** Multiple cores trying to access the same cache lines or evicting each other's data (Cache Pollution).
- **Thread Overhead:** For smaller  $N$  (or efficient  $N$ ), the cost of spawning and syncing threads outweighs the compute benefits (evident in the spikes).

## Bottleneck Shift

**Naive:** CPU Waiting for RAM.



**OpenMP:** *Multiple CPUs Fighting for the Same RAM Channel.*