


# **Optimization Strategies: Loop Unrolling & Register Blocking**

## **SIMD 1D vs SIMD 2D (Register Blocking)**

Progetto AMSC


November 27, 2025

# Part 1: SIMD 1D Unrolling Overview



**Figure: Execution Time (1D)**

*Linear scaling, but limited by memory bandwidth.*



**Figure: Throughput (1D)**

*Performance caps at  $\approx 15$  GFLOPS for large  $N$ .*

# 1D Unrolling: Quantitative Analysis

## 1D Unrolling Performance (Float vs Double)


| Size ( $N$ ) | Float         | Double        | Gap                |
|--------------|---------------|---------------|--------------------|
| $128^3$      | 0.16 ms       | 0.30 ms       | +80%               |
| $1024^3$     | 75.87 ms      | 160.20 ms     | +111%              |
| $1536^3$     | 255.82 ms     | 1.07 s        | +317%              |
| $2048^3$     | <b>1.10 s</b> | <b>3.00 s</b> | <b>2.7x Slower</b> |

### Observation

1D Unrolling reduces loop overhead compared to basic SIMD, but it hits the **Memory Wall** early.


At  $N = 2048$ , the CPU spends more time fetching matrix  $B$  from RAM than computing, limiting the Float throughput to  $\approx 15.6$  GFLOPS.

## Part 2: SIMD 2D Unrolling (Register Blocking)



**Figure: Execution Time (2D)**

*Significant speedup observed at large N.*



**Figure: Throughput (2D)**

*Higher sustained throughput ( $\approx 23.5$  GFLOPS).*

# 2D Unrolling: Quantitative Analysis

## 2D Unrolling Performance (Float vs Double)

| Size ( $N$ ) | Float         | Double        | Gap                |
|--------------|---------------|---------------|--------------------|
| $128^3$      | 0.13 ms       | 0.26 ms       | +100%              |
| $1024^3$     | 68.68 ms      | 137.71 ms     | +100%              |
| $2048^3$     | <b>0.73 s</b> | <b>1.86 s</b> | <b>2.5x Slower</b> |

Speedup vs 1D (Float):  $1.10s \rightarrow 0.73s$  (+50% Faster)

### The Impact of Register Blocking

By processing blocks of data, we improve **Arithmetic Intensity**.

We achieve  $\approx 23.5$  GFLOPS (Float) vs 15.6 GFLOPS (1D). The CPU is utilized much more efficiently because it waits less for memory.

# Optimization Strategy: 1D Unrolling (Horizontal)

**Mechanism:** For each scalar element  $A_{i,k}$ , we load a "horizontal strip" of  $B$  and update a corresponding strip of  $C$ .

- **Loads:** 1 Scalar ( $A$ ) +  $N$  Vectors ( $B$ ).
- **Compute:** Update  $N$  accumulators in  $C$  (same row).

## Pros & Cons:

- + **Easy to implement:** Only requires unrolling the innermost loop ( $j$ ).
- + **Low Register Pressure:** Uses few registers (1 for  $A$ , 4-8 for  $C$ ).
- **Bandwidth Waste:** For every new row of  $C$ , we must reload the entire matrix  $B$ . The CPU spends cycles waiting for loads.

## 1D Pattern

$$C[i][j \dots j+4] += A[i][k] \times B[k][j \dots j+4]$$

# Optimization Strategy: 2D Register Blocking (Pro Level)

**Mechanism:** We process a block of rows (e.g., 2 rows of  $A$ ) simultaneously. We reuse the *same* vector of  $B$  to update multiple rows of  $C$ .

- **Loads:** 2 Scalars ( $A_i, A_{i+1}$ ) + **1 Vector (B)**.
- **Compute:** 2 FMA operations simultaneously:

$$C_i += A_i \times B$$

$$C_{i+1} += A_{i+1} \times B$$

## Why is it faster?

- + **Halved Memory Access:** We fetch  $B$  once but use it for 2 rows of  $C$ . This drastically reduces L1 Cache traffic.
- + **Higher Arithmetic Intensity:** More math per byte loaded.

## Trade-offs

- ! **High Register Pressure:** Requires many active registers. Excessive blocking (e.g., 4x4) can cause *Register Spilling*, degrading performance.
- ! **Complexity:** Requires handling edge cases (odd matrix sizes) with specific cleanup loops.